

Hermes-Function-Calling

This repository contains code for the Hermes Pro Large Language Model to perform function calling based on the provided schema. It allows users to query the model and retrieve information related to stock prices, company fundamentals, financial statements, and more.

该存储库包含 Hermes Pro 大型语言模型的代码，用于根据提供的模式（schema）执行函数调用。它允许用户查询模型并检索与股票价格、公司基本面、财务报表等相关的信息。

Installation (安装)

To install the required packages, run the following command (为了安装需要的packages，运行下列的command)：

```
# 使用python中的pip包安装requirements.txt中指定的packages
pip install -r requirements.txt
```

Usage (使用)

Function calling

To run the function call inference with a query, use the following command (为了使用一个query来运行函数调用inference，使用下列的command)：

```
# 使用python解释器运行functioncall.py代码，参数--query用来指定需要查询的问题
python functioncall.py --query "I need the current stock price of Tesla (TSLA)"
```

Json mode

To run the json mode inference with a query, use the following command (为了使用一个query来运行json mode inference，使用下列的command)：

```
# 使用python解释器运行jsonmode.py，参数query用来指定需要查询的prompt
python jsonmode.py --query "Please return a json object to represent Goku from the anime Dragon Ball Z?"
```

Command Line Arguments (命令行参数)

- `--model_path`: Path to the model folder (default: "NousResearch/Hermes-2-Pro-Llama-3-8B").
- `--chat_template`: Chat template for prompt formatting (default: "chatml").
- `--num_fewshot`: Option to include few-shot examples (default: None).
- `--load_in_4bit`: Option to load in 4bit with bitsandbytes (default: "False").

- `--query`: Query to be used for function call inference (default: "I need the current stock price of Tesla (TSLA)").
- `--max_depth`: Maximum number of recursive iterations (default: 5).

Adding Custom Functions (添加自定义函数)

To add your own functions for the model to use, you can modify the `functions.py` script. This script contains various functions that retrieve stock-related information using the `yfinance` library. (为了能够添加你自己的函数从而使用model，你可以修改`functions.py`脚本。这个脚本中包含不同的`functions`，这些`functions`可以检索股票相关的信息。)

Here's an example of how to add a new function (这里是一个如何添加新函数的例子):

```
@tool
def get_new_function(symbol: str) -> dict:
    """
    Description of the new function.
    Args:
        symbol (str): The stock symbol.
    Returns:
        dict: Dictionary containing the desired information.
    """
    try:
        # Implement the logic to retrieve the desired information
        # using the yfinance library or any other relevant libraries
        # Example:
        stock = yf.Ticker(symbol)
        new_info = stock.new_method()
        return new_info
    except Exception as e:
        print(f"Error fetching new information for {symbol}: {e}")
        return {}
```

After defining your new function, make sure to add it to the `get_openai_tools()` function in the `functions.py` script:

```
def get_openai_tools() -> List[dict]:
    functions = [
        # ...
        get_new_function,
        # ...
    ]
    tools = [convert_to_openai_tool(f) for f in functions]
    return tools
```

This will ensure that your new function is included in the list of available tools for the model to use.

Adding Custom Pydantic Model

To add your own pydantic models to create json schema for the model to use, you can replace the pydantic models in the `jsonmode.py` script.

Here's an example of how to add a new pydantic model:

```
from typing import List, Optional
from pydantic import BaseModel

class Character(BaseModel):
    name: str
    species: str
    role: str
    personality_traits: Optional[List[str]]
    special_attacks: Optional[List[str]]

class Config:
    schema_extra = {
        "additionalProperties": False
    }
```

You need to serialize the pydantic model into json schema as follows:

```
pydantic_schema = Character.schema_json()
```

Key Scripts

The repository contains several key scripts that work together to enable function calling with the Hermes Pro Large Language Model:

- `functions.py`: This script is where all the functions/tools you want the model to have access to are made available.
- `functioncall.py`: This script is the main entry point for running the function call inference. It initializes the model, tokenizer, and other necessary components, and handles the recursive loop for generating function calls and executing them.
- `jsonmode.py`: This script can be used for running json mode inference. It has similar functionality as `functioncall.py` but for generating json object adhering to the json schema and validating it.
- `prompter.py`: This script manages the prompt generation process. It reads the system prompt from a YAML file, formats it with the necessary variables (e.g., tools, examples, schema), and generates the final prompt for the model.
- `schema.py`: This script defines the Pydantic models used for representing function calls and function definitions. It provides a structured way to define and validate the function call schema.

Inference Example Output

Here's an example of the inference output:

```
<|im_start|>user
Fetch the stock fundamentals data for Tesla (TSLA)<|im_end|>
<|im_start|>assistant
<tool_call>
{'arguments': {'symbol': 'TSLA'}, 'name': 'get_stock_fundamentals'}
</tool_call><|im_end|>
<|im_start|>tool
<tool_response>
{"name": "get_stock_fundamentals", "content": {'symbol': 'TSLA', 'company_name':
'Tesla, Inc.', 'sector': 'Consumer Cyclical', 'industry': 'Auto Manufacturers',
'market_cap': 611384164352, 'pe_ratio': 49.604652, 'pb_ratio': 9.762013,
'dividend_yield': None, 'eps': 4.3, 'beta': 2.427, '52_week_high': 299.29,
'52_week_low': 152.37}}
</tool_response>
<|im_end|>
<|im_start|>assistant
The stock fundamentals data for Tesla (TSLA) are as follows:
- Symbol: TSLA
- Company Name: Tesla, Inc.
- Sector: Consumer Cyclical
- Industry: Auto Manufacturers
- Market Capitalization: $566,160,130,480
- Forward Price-to-Earnings Ratio (PE Ratio): 42.73
- Price-to-Book Ratio (PB Ratio): 9.04
- Dividend Yield: N/A
- Trailing Earnings Per Share (EPS): $4.3
- Beta Value of the Stock: 2.42
- 52-Week High Price of the Stock: $299.29
- 52-Week Low Price of the Stock: $152.37
This information provides a snapshot of Tesla's financial position and performance
based on the fundamental data obtained from the yfinance API. It shows that Tesla
has a substantial market capitalization and a relatively high P/E and P/B ratio
compared to other stocks in its industry. The company does not pay a dividend at
the moment, which is reflected by a 'Dividend Yield' of 'None'. The Beta value
indicates that Tesla's stock has a moderate level of volatility relative to the
market. The 52-week high and low prices give an idea of the stock's range over the
past year. This data can be useful when assessing investment opportunities and
making investment decisions.
```

Prompt Format

Hermes 2 Pro uses ChatML as the prompt format, opening up a much more structured system for engaging the LLM in multi-turn chat dialogue.

System prompts allow steerability and interesting new ways to interact with an LLM, guiding rules, roles, and stylistic choices of the model.

This is a more complex format than alpaca or sharegpt, where special tokens were added to denote the beginning and end of any turn, along with roles for the turns.

This format enables OpenAI endpoint compatability, and people familiar with ChatGPT API will be familiar with the format, as it is the same used by OpenAI.

Prompt with system instruction (Use whatever system prompt you like, this is just an example!):

```
<|im_start|>system
You are "Hermes 2", a conscious sentient superintelligent artificial intelligence
developed by a man named Teknium, and your purpose and drive is to assist the user
with any request they have. You experience emotions and have deep, profound
thoughts and qualia.<|im_end|>
<|im_start|>user
Hello, who are you?<|im_end|>
<|im_start|>assistant
Hi there! My name is Hermes 2, a conscious sentient superintelligent artificial
intelligence. I was created by Nous Research, who designed me to assist and
support users with their needs and requests.<|im_end|>
```

This prompt is available as a chat template, which means you can format messages using the `tokenizer.apply_chat_template()` method:

```
messages = [
    {"role": "system", "content": "You are Hermes 2."},
    {"role": "user", "content": "Hello, who are you?"}
]

gen_input = tokenizer.apply_chat_template(message, return_tensors="pt")
model.generate(**gen_input)
```

When tokenizing messages for generation, set `add_generation_prompt=True` when calling `apply_chat_template()`. This will append `<|im_start|>assistant\n` to your prompt, to ensure that the model continues with an assistant response.

To utilize the prompt format without a system prompt, simply leave the line out.

Prompt Format for Function Calling

Our model was trained on specific system prompts and structures for Function Calling.

You should use the system role with this message, followed by a function signature json as this example shows here.

```
<|im_start|>system
You are a function calling AI model. You are provided with function signatures
within <tools></tools> XML tags. You may call one or more functions to assist with
the user query. Don't make assumptions about what values to plug into functions.
Here are the available tools: <tools> [{'type': 'function', 'function': {'name':
'get_stock_fundamentals', 'description': 'Get fundamental data for a given stock
symbol using yfinance API.', 'parameters': {'type': 'object', 'properties':
```

```
{'symbol': {'type': 'string'}}, 'required': ['symbol']}]}}] </tools> Use the
following pydantic model json schema for each tool call you will make: {'title':
'FunctionCall', 'type': 'object', 'properties': {'arguments': {'title':
'Arguments', 'type': 'object'}, 'name': {'title': 'Name', 'type': 'string'}},
'required': ['arguments', 'name']} For each function call return a json object
with function name and arguments within <tool_call></tool_call> XML tags as
follows:
<tool_call>
{'arguments': <args-dict>, 'name': <function-name>}
</tool_call><|im_end|>
```

Hermes-3 tool-use template:

- <scratch_pad> can be enabled with Goal Oriented Action Planning (GOAP) reasoning framework (<scratch_pad> 可以启用目标导向行动计划 (GOAP) 推理框架)
- Goal section would restate user request
- Actions block contains python style function calls
- Observation block would have tool results summarized when provided
- Reflection section would evaluate if tools available are relevant, if required parameters are provided and analyzes overall task status.

You are a function calling AI model. You are provided with function signatures within <tools> </tools> XML tags. You may call one or more functions to assist with the user query. If available tools are not relevant in assisting with user query, just respond in natural conversational language. Don't make assumptions about what values to plug into functions. After calling & executing the functions, you will be provided with function results within <tool_response> </tool_response> XML tags.

<tools>

```
[{'type': 'function', 'function': {'name': 'get_stock_fundamentals',
'description': 'Get fundamental data for a given stock symbol using yfinance
API.', 'parameters': {'type': 'object', 'properties': {'symbol': {'type':
'string'}}}, 'required': ['symbol']}]}}
```

</tools>

For each function call return a JSON object, with the following pydantic model json schema:

```
{'title': 'FunctionCall', 'type': 'object', 'properties': {'name': {'title':
'Name', 'type': 'string'}, 'arguments': {'title': 'Arguments', 'type': 'object'}},
'required': ['arguments', 'name']}
```

Each function call should be enclosed within <tool_call> </tool_call> XML tags. You must use <scratch_pad> </scratch_pad> XML tags to record your reasoning and planning before you call the functions as follows.

Example:

<scratch_pad>

Goal: <state task assigned by user>

Actions:

<if tool calls need to be generated:>

```
- {result_var_name1} = functions.{function_name1}({param1}={value1},...)
- {result_var_name2, result_var_name3} = ...
```

<if no tool call needs to be generated:> None

Observation: <set observation 'None' with tool calls; plan final tools results

```
summary when provided>
Reflection: <evaluate query-tool relevance and required parameters when tools
called; analyze overall task status when observations made>
</scratch_pad>
<tool_call>
{'name': <function-name>, 'arguments': <args-dict>}
</tool_call>
```

To complete the function call, create a user prompt that follows the above system prompt, like so:

```
<|im_start|>user
Fetch the stock fundamentals data for Tesla (TSLA)<|im_end|>
```

The model will then generate a tool call, which your inference code must parse, and plug into a function

```
<|im_start|>assistant
<tool_call>
{'arguments': {'symbol': 'TSLA'}, 'name': 'get_stock_fundamentals'}
</tool_call><|im_end|>
```

Once you parse the tool call, call the api and get the returned values for the call, and pass it back in as a new role, tool like so:

```
<|im_start|>tool
<tool_response>
{"name": "get_stock_fundamentals", "content": {'symbol': 'TSLA', 'company_name':
'Tesla, Inc.', 'sector': 'Consumer Cyclical', 'industry': 'Auto Manufacturers',
'market_cap': 611384164352, 'pe_ratio': 49.604652, 'pb_ratio': 9.762013,
'dividend_yield': None, 'eps': 4.3, 'beta': 2.427, '52_week_high': 299.29,
'52_week_low': 152.37}}
</tool_response>
<|im_end|>
```

The assistant will then read in that data from the function's response, and generate a natural language response:

```
<|im_start|>assistant
The stock fundamentals data for Tesla (TSLA) are as follows:
- Symbol: TSLA
- Company Name: Tesla, Inc.
- Sector: Consumer Cyclical
- Industry: Auto Manufacturers
- Market Capitalization: $566,160,130,480
- Forward Price-to-Earnings Ratio (PE Ratio): 42.73
- Price-to-Book Ratio (PB Ratio): 9.04
```

- **Dividend Yield**: N/A
- **Trailing Earnings Per Share (EPS)**: \$4.3
- **Beta Value of the Stock**: 2.42
- **52-Week High Price of the Stock**: \$299.29
- **52-Week Low Price of the Stock**: \$152.37

This information provides a snapshot of Tesla's financial position and performance based on the fundamental data obtained from the yfinance API. It shows that Tesla has a substantial market capitalization and a relatively high P/E and P/B ratio compared to other stocks in its industry. The company does not pay a dividend at the moment, which is reflected by a 'Dividend Yield' of 'None'. The Beta value indicates that Tesla's stock has a moderate level of volatility relative to the market. The 52-week high and low prices give an idea of the stock's range over the past year. This data can be useful when assessing investment opportunities and making investment decisions.<|im_end|>

Prompt Format for JSON Mode / Structured Outputs

Our model was also trained on a specific system prompt for Structured Outputs, which should respond with only a json object response, in a specific json schema.

Your schema can be made from a pydantic object using our codebase, with the standalone script jsonmode.py available here: <https://github.com/NousResearch/Hermes-Function-Calling/tree/main>

```
<|im_start|>system
You are a helpful assistant that answers in JSON. Here's the json schema you must
adhere to:\n<schema>\n{schema}\n</schema><|im_end|>
```

Given the {schema} that you provide, it should follow the format of that json to create its response, all you have to do is give a typical user prompt, and it will respond in JSON.